

Regression

Liang Liang

Regression

- Regression is a subcategory of supervised learning where the goal is to predict continuous target value given the features of a sample.
- The relationship between the features and the target value may be linear or nonlinear.
- The target value could be a scalar or a vector.

Linear Regression

$\hat{y} = w_{(1)}x_{(1)} + w_{(2)}x_{(2)} + b$, a linear model

$$y = \hat{y} + \varepsilon$$

a data point x is a feature vector $[x_{(1)}, x_{(2)}]$

y is the 'true' target value (ground-truth)

\hat{y} is the predicted target value from the model

ε is something that can not be explained by the linear model

ε is treated as 'random noise'

$\{w_{(1)}, w_{(2)}, b\}$ are the parameters of the linear model

Linear Regression

simple linear regression: $\hat{y} = w_{(1)}x + b$

b is intercept

w is slope

x is a scalar

multiple linear regression: $\hat{y} = w_{(1)}x_{(1)} + w_{(2)}x_{(2)} \dots + w_{(M)}x_{(M)} + b$

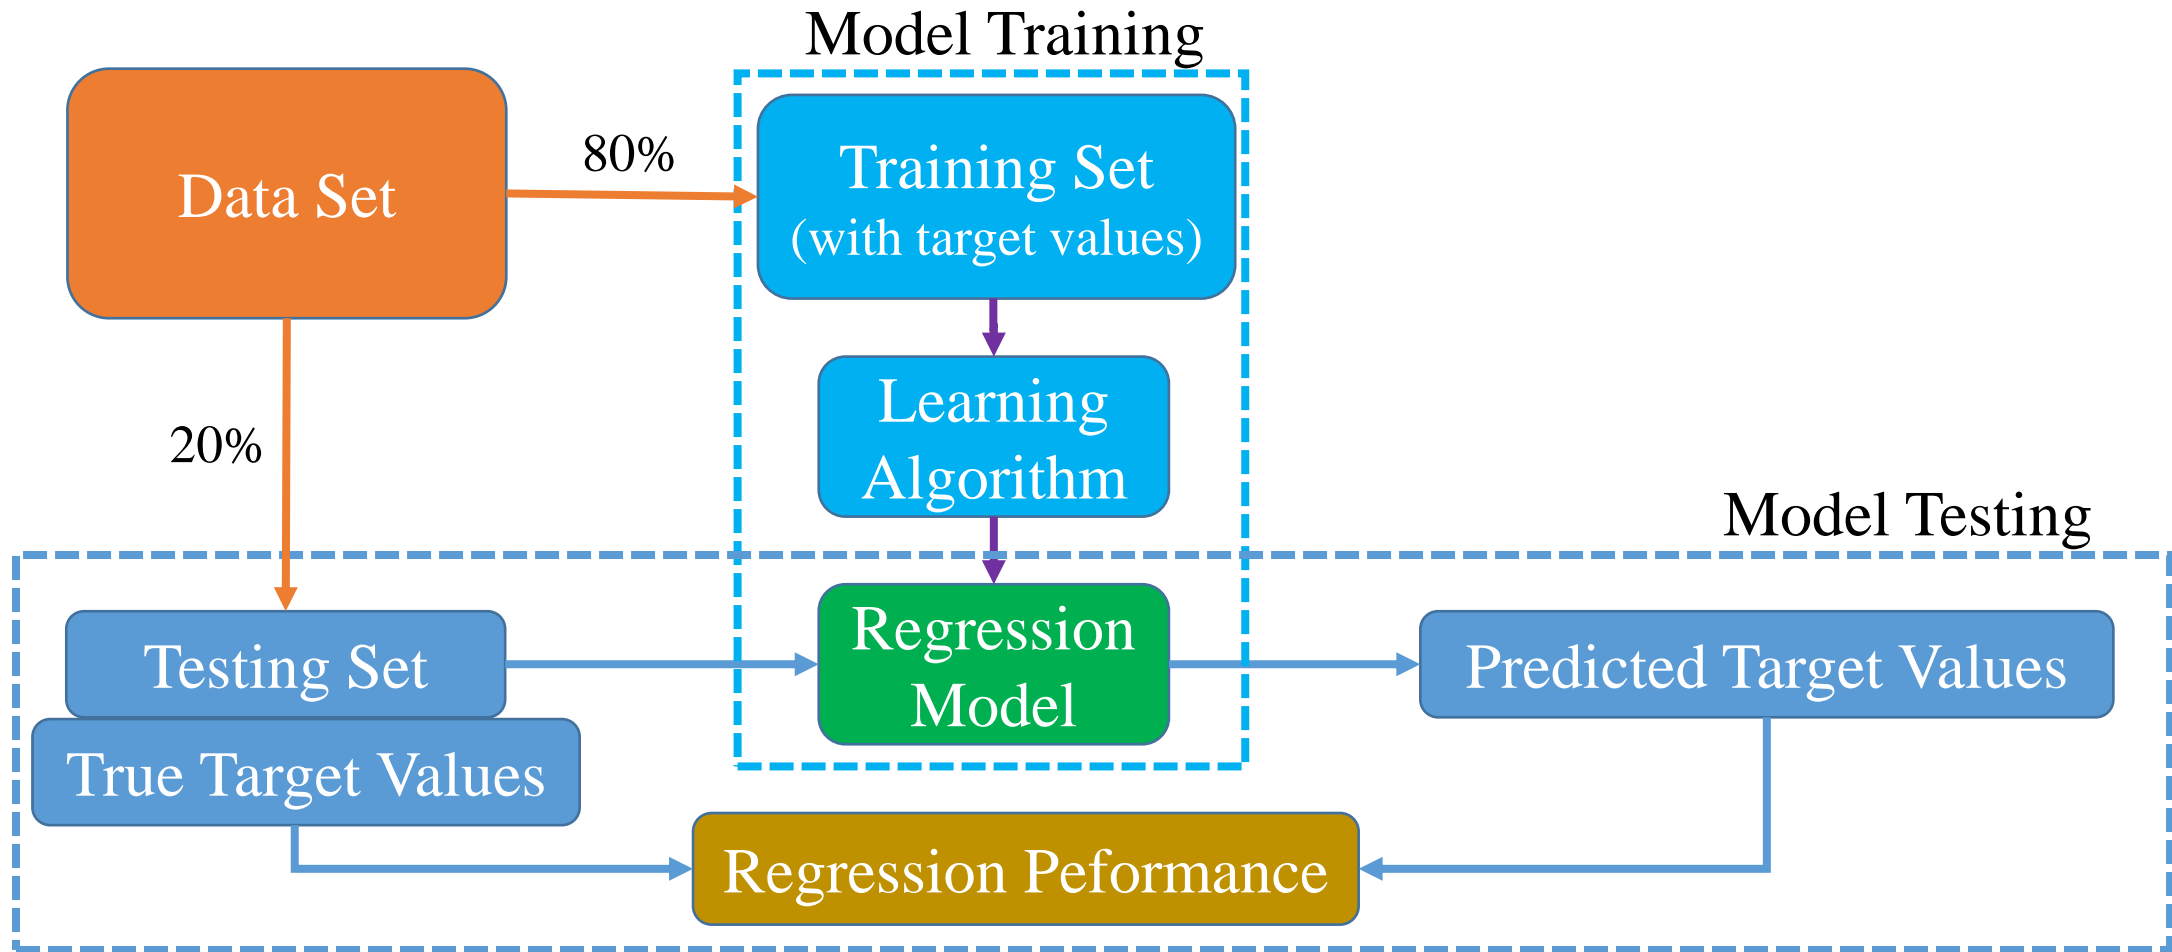
$\{w_{(1)}, w_{(2)}, \dots, w_{(M)}, b\}$ are the parameters of the linear model

$x = [x_{(1)}, x_{(2)}, \dots, x_{(M)}]$ is a sample (feature vector)

Fit the linear model to training dataset, to obtain the optimal parameters

Evaluate the trained/fitted linear model on testing dataset

the workflow of a regression study

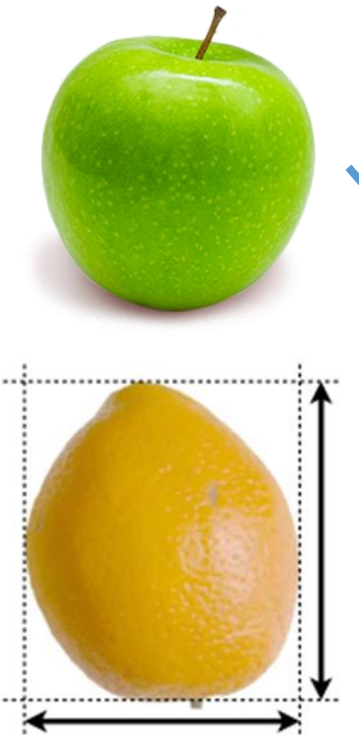


Example: linear regression on the fruit dataset

A bucket
of fruits

The fruit dataset was created by Dr. Iain Murray at the University of Edinburgh. He bought a few dozen oranges, lemons and apples, and recorded their features in a table.

The fruit dataset { 1:apple, 2:mandarin, 3:orange, 4:lemon },
Each row contains the information of a fruit sample/instance



The image shows a green apple and a yellow lemon. The apple is at the top left, and the lemon is at the bottom left. A blue arrow points from the apple to the first row of the table, and another blue arrow points from the lemon to the second row. The lemon has black arrows indicating its width and height, with dashed lines forming a bounding box.

fruit label	fruit_name	subtype	mass (g)	width (cm)	height (cm)	color_score
1	apple	granny_smith	192	8.4	7.3	0.55
4	lemon	spanish_belsan	194	7.2	10.3	0.70

Example: linear regression on the fruit dataset

first step: load the dataset

The table has 59 rows (samples)

```
1 fruits = pd.read_table('fruit_data_with_colors.txt')
```

```
1 fruits
```

	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79
5	2	mandarin	mandarin	80	5.8	4.3	0.77
6	2	mandarin	mandarin	80	5.9	4.3	0.81
7	2	mandarin	mandarin	76	5.8	4.0	0.81
8	1	apple	braeburn	178	7.1	7.8	0.92
9	1	apple	braeburn	172	7.4	7.0	0.89

linear regression on the fruit dataset

goal: predict mass given width, height and color

fruit label	fruit_name	subtype	mass (g)	width (cm)	height (cm)	color_score
1	apple	granny_smith	192	8.4	7.3	0.55

Feature Vector $x_n = [x_{(n,1)}, x_{(n,2)}, x_{(n,3)}]$



$x_{(n,1)}$: width
 $x_{(n,2)}$: height
 $x_{(n,3)}$: color



Linear
Model



Target

mass \hat{y}_n

n is the index of the sample
(row index in the table)

$$\hat{y}_n = w_{(1)}x_{(n,1)} + w_{(2)}x_{(n,2)} + w_{(3)}x_{(n,3)} + b$$

linear regression on the fruit dataset

- data splitting

split the data (59 samples) into a training dataset (80%) and a testing dataset (20%)

```
1 feature_names = ['width', 'height', 'color_score']  
2 feature_names
```

```
['width', 'height', 'color_score']
```

```
1 target_name = ['mass']  
2 target_name
```

```
['mass']
```

Split the data into a Training dataset and a Testing dataset

```
1 X = fruits[feature_names]  
2 Y = fruits[target_name]  
3 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=0)
```

Model Training: fit the model to the training dataset

The training set contains N input-output pairs: $\{(x_n, y_n), n = 1, \dots, N\}$

x_n is a feature vector ; y_n is the true target value

MSE (mean squared error) loss function:

$$L(w_{(1)}, w_{(2)}, w_{(3)}, b) = \frac{1}{N} \sum_{n=1}^N (y_n - \hat{y}_n)^2$$

$$\hat{y}_n = w_{(1)}x_{(n,1)} + w_{(2)}x_{(n,2)} + w_{(3)}x_{(n,3)} + b$$

The goal of training is to find the best parameters $\{w_{(1)}, w_{(2)}, w_{(3)}, b\}$ such that the loss function is minimized. After training, the prediction \hat{y}_n from the model should be very close to the true target y_n

The loss function

- Given N training data points with target values $\{(x_n, y_n), n = 1, \dots, N\}$,
 - $x_n = [x_{(n,1)}, x_{(n,2)}, \dots, x_{(n,M)}]$ is a data sample / point / feature vector
 - M is the number of components/attributes, and it is 3 for the fruit dataset
 - N is the number of samples
- find the best parameters that minimize the loss:

$$L(w_{(1)}, \dots, w_{(M)}, b) = \frac{1}{N} \sum_{n=1}^N (y_n - \hat{y}_n)^2$$

$$\hat{y}_n = w_{(1)}x_{(n,1)} + w_{(2)}x_{(n,2)} + \dots + w_{(M)}x_{(n,M)} + b$$

$$\hat{y}_n = w_{(1)}x_{(n,1)} + w_{(2)}x_{(n,2)} + \cdots + w_{(M)}x_{(n,M)} + b$$

$$x_n = [x_{(n,1)}, x_{(n,2)}, \dots, x_{(n,M)}]$$

$$\text{Let } w = [w_{(1)}, \dots, w_{(M)}]$$

$$\text{Let } w \cdot x_n = w_{(1)}x_{(n,1)} + w_{(2)}x_{(n,2)} + \cdots + w_{(M)}x_{(n,M)}$$

then we have

$$\hat{y}_n = w \cdot x_n + b$$

$$L(w, b) = \frac{1}{N} \sum_{n=1}^N (y_n - \hat{y}_n)^2 \quad \text{and} \quad \hat{y}_n = w \cdot x_n + b$$

$$x_n = [x_{(n,1)}, x_{(n,2)}, \dots, x_{(n,M)}] \quad \text{and} \quad w = [w_{(1)}, \dots, w_{(M)}]$$

$$\text{Let } \frac{\partial L}{\partial w} = \left[\frac{\partial L}{\partial w_{(1)}}, \frac{\partial L}{\partial w_{(2)}}, \dots, \frac{\partial L}{\partial w_{(M)}} \right]$$

Then we obtain:

$$\frac{\partial L}{\partial w} = -\frac{2}{N} \sum_{n=1}^N (y_n - w \cdot x_n - b) \times x_n$$

- Use gradient descent to find the optimal parameters w and b

step-0: initialize w and b randomly

step-1: compute $\frac{\partial L}{\partial w} = -\frac{2}{N} \sum_{n=1}^N (y_n - w \cdot x_n - b) \times x_n$

$$\frac{\partial L}{\partial b} = -\frac{2}{N} \sum_{n=1}^N (y_n - w \cdot x_n - b)$$

step-2: update $w \leftarrow w - \eta \frac{\partial L}{\partial w}$ and $b \leftarrow b - \eta \frac{\partial L}{\partial b}$

where η is learning rate

repeat step-1 and step-2 for a large number of iterations

Model Testing: apply the model to the testing dataset

The testing set contains K input-output pairs: $\{(x_k, y_k), k = 1, \dots, K\}$

mean squared error (MSE)

$$MSE = \frac{1}{K} \sum_{k=1}^K (y_k - \hat{y}_k)^2$$

mean absolute error (MAE)

$$MAE = \frac{1}{K} \sum_{k=1}^K |y_k - \hat{y}_k|$$

Mean absolute percentage error (MAPE)

$$MAPE = \frac{1}{K} \sum_{k=1}^K \left| \frac{y_k - \hat{y}_k}{y_k} \right| \times 100\%$$

```
#prediction on the testing dataset
Y_test_pred = linear_model.predict(X_test)
MSE = np.mean((Y_test - Y_test_pred)**2)
MAE = np.mean(np.abs(Y_test - Y_test_pred))
MAPE = np.mean(np.abs(Y_test - Y_test_pred)/Y_test)
print('MSE=', MSE)
print('MAE=', MAE)
print('MAPE=', MAPE)
```

we do not need for loops, use vectorized operation

linear_regression_on_fruit_data.ipynb

LinearRegression_implementation.ipynb

Nonlinear regression using a polynomial model

- A polynomial model of degree K

$$\hat{y} = w_{(1)}x + w_{(2)}x^2 \dots + w_{(K)}x^K + b$$

where x is a scalar

- Convert the polynomial model to a linear model:

$$\hat{y} = w_{(1)}\tilde{x}_{(1)} + w_{(2)}\tilde{x}_{(2)} + \dots + w_{(K)}\tilde{x}_{(K)} + b$$

$\tilde{x} = [\tilde{x}_{(1)}, \tilde{x}_{(2)}, \dots, \tilde{x}_{(K)}] = [x, x^2, \dots, x^K]$ is a feature vector

- The optimal values of the parameters $\{w_{(1)}, \dots, w_{(K)}, b\}$ can be obtained by using the gradient descent method with the MSE loss.